

Parallelizing Loops

Moreno Marzolla
Dip. di Informatica—Scienza e Ingegneria (DISI)
Università di Bologna

Copyright © 2017–2026
Moreno Marzolla, Università di Bologna, Italy
<https://www.moreno.marzolla.name/teaching/CP/>



Except where stated otherwise, this work is licensed under the Creative Commons Attribution-ShareAlike 4.0 International License (CC BY-SA 4.0). To view a copy of this license, visit <https://creativecommons.org/licenses/by-sa/4.0/> or send a letter to Creative Commons, 543 Howard Street, 5th Floor, San Francisco, California, 94105, USA.

Credits

- Salvatore Orlando (Univ. Ca' Foscari di Venezia)
- Mary Hall (Univ. of Utah)


Loop optimization

- 90% of execution time in 10% of the code
 - Mostly in loops.
- Loop optimizations
 - Transform loops preserving the semantics.
- Goal:
 - Single-threaded system: mostly optimizing for memory hierarchy.
 - Multi-threaded and vector systems: **loop parallelization.**

Executing loops in parallel

```
for (i=0; i<N; i++) {  
    foo(i);  
}
```

*Each green box is
an independent
execution unit*



- In some cases loop iterations are **independent** and can be executed **in parallel**
 - The number of execution units could be much less than the number of loop iterations.
 - Therefore, each execution unit takes care of multiple loop iterations in some **unspecified order**.


Data Dependence

- Two memory accesses are involved in a **data dependence** if they may refer to the same memory location and one of the references is a write.

- Data-Flow or true dependence**

- RAW (Read After Write).


S1	<code>a = b + c;</code>
S2	<code>d = 2 * a;</code>



- Anti dependence**

- WAR (Write After Read).


S1	<code>c = a + b;</code>
S2	<code>a = 2 * a;</code>



- Output dependence**

- WAW (Write After Write).


S1	<code>a = k;</code>
S2	<code>if (a>0) {</code>
	<code> a = 2 * c;</code>
	<code>}</code>



Control Dependence


- An instruction S2 has a **control dependence** on S1 if the outcome of S1 determines whether S2 is executed or not.
 - Of course, S1 and S2 can not be exchanged.
- This type of dependence applies to the condition of an if-then-else or loop with respect to their bodies.

```
S1 if (a > 0) {  
S2     a = 2 * c;  
     } else {  
S3     b = 3;  
     }
```



Dependence

- In the following, we use a simple arrow to denote any type of dependence.
 - S2 depends on S1.




Fundamental Theorem of Dependence

- *“Any reordering transformation that preserves every dependence in a program preserves the meaning of that program”*
- Recognizing parallel loops (intuitively):
 - Find data dependences in a loop.
 - No dependences crossing iteration boundary → loop can be parallelized.

Example


```
for (i=0; i<n; i++) {  
  S1 a[i] = b[i] + c[i];  
}
```



- Each iteration **does not depend** on previous ones.
 - There are no dependences crossing iteration boundaries.
- This loop is **fully parallelizable**.
 - Loop iterations can be performed **in parallel**, in any order.
 - Iterations can be split across processors.

Example


```
for (i=1; i<n; i++) {  
  S1 a[i] = a[i-1] + b[i]  
}
```



- Each iteration **depends** on the previous one (RAW).
 - The value **a[i-1]** read at iteration **i** has been modified at iteration **i-1**;
 - **Loop-carried dependence**.
- Hence, this loop is **not** parallelizable with **trivial** transformations.

Example


```
s = 0;  
for (i=0; i<n; i++) {  
  S1 s = s + a[i];  
}
```



- We have a **loop-carried dependence** on **s** that can not be removed with trivial loop transformations...
- ...but can be removed with *non-trivial* transformations.
 - This is a **reduction**, so we can apply the reduction parallel pattern.


Exercise: draw the dependences among iterations of the following loop

```
for (i=2; i<n; i++) {  
  S1 a[i] = 4 * c[i-1] - 2;  
  S2 b[i] = a[i] * 2;  
  S3 c[i] = a[i-1] + 3;  
  S4 d[i] = b[i] + c[i-2];  
}
```




Exercise: draw the dependences among iterations of the following loop

```
for (i=2; i<n; i++) {  
  S1 a[i] = 4 * c[i-1] - 2;  
  S2 b[i] = a[i] * 2;  
  S3 c[i] = a[i-1] + 3;  
  S4 d[i] = b[i] + c[i-2];  
}
```




Exercise: draw the dependences among iterations of the following loop

```
for (i=2; i<n; i++) {  
  S1 a[i] = 4 * c[i-1] - 2;  
  S2 b[i] = a[i] * 2;  
  S3 c[i] = a[i-1] + 3;  
  S4 d[i] = b[i] + c[i-2];  
}
```



Exercise: draw the dependences among iterations of the following loop


```
for (i=2; i<n; i++) {  
  S1 a[i] = 4 * c[i-1] - 2;  
  S2 b[i] = a[i] * 2;  
  S3 c[i] = a[i-1] + 3;  
  S4 d[i] = b[i] + c[i-2];  
}
```



Parallelizing Loops

Exercise: draw the dependences among iterations of the following loop


```
for (i=2; i<n; i++) {  
  S1 a[i] = 4 * c[i-1] - 2;  
  S2 b[i] = a[i] * 2;  
  S3 c[i] = a[i-1] + 3;  
  S4 d[i] = b[i] + c[i-2];  
}
```



Parameterizing Loops

Exercise: draw the dependences among iterations of the following loop


```
for (i=2; i<n; i++) {  
  S1 a[i] = 4 * c[i-1] - 2;  
  S2 b[i] = a[i] * 2;  
  S3 c[i] = a[i-1] + 3;  
  S4 d[i] = b[i] + c[i-2];  
}
```



Parameterizing Loops

Exercise: draw the dependences among iterations of the following loop

```
for (i=2; i<n; i++) {  
  S1 a[i] = 4 * c[i-1] - 2;  
  S2 b[i] = a[i] * 2;  
  S3 c[i] = a[i-1] + 3;  
  S4 d[i] = b[i] + c[i-2];  
}
```




Parallelizing Loops


Removing dependences: Aligning

- Dependences can sometimes be removed by **aligning** loop iterations.

```
a[0] = 0;  
for (i=1; i<n; i++) {  
  S1 a[i] = b[i-1] * c[i];  
  S2 d[i] = a[i-1] + 2;  
}
```




```
a[0] = 0;  
d[1] = a[0] + 2;  
for (i=1; i<n-1; i++) {  
  T1 a[i] = b[i-1] * c[i];  
  T2 d[i+1] = a[i] + 2;  
}  
a[n-1] = b[n-2] * c[n-1];
```



RAW


*a[i-1] is read at iteration i
and written at iteration i-1.*



Removing dependences: Fission


- In this example, dependences can be removed by using a temporary array and splitting the loop.

```
for (i=n-1; i>0; i--) {  
  S1 a[i] = a[i-1];  
}
```



Write-After-Read (WAR)
 $a[i]$ is written at iteration i and read at iteration $i+1$.

```
b[0] = a[0];  
for (i=1; i<n; i++) {  
  T1 b[i] = a[i];  
  T2 a[i] = b[i-1];  
}
```




RAW
 $b[i-1]$ is read at iteration i and written at iteration $i-1$.

WAR
 $a[i]$ is written at iteration i after it is read.


Removing dependences: Fission

- In this example, dependences can be removed by using a temporary array and splitting the loop.

```
for (i=n-1; i>0; i--) {  
S1  a[i] = a[i-1];  
}
```



```
b[0] = a[0];  
for (i=1; i<n; i++) {  
T1  b[i] = a[i];  
T2  a[i] = b[i-1];  
}
```



```
for (i=0; i<n; i++) {  
    b[i] = a[i];  
}  
for (i=1; i<n; i++) {  
    a[i] = b[i-1];  
}
```


*Both loops can
be parallelized*

Optimizing loops: Fusion

- Sometimes it is possible to merge multiple **parallelizable** loops to reduce overhead.

```
for (i=0; i<n; i++) {  
    a[i] = b[i] * c[i];  
}  
for (i=0; i<n; i++) {  
    b[i] = f(d[i]) * h[i];  
}
```

Both loops are parallelizable, since they have no loop-carried dependencies.





```
for (i=0; i<n; i++) {  
    a[i] = b[i] * c[i];  
    b[i] = f(d[i]) * h[i];  
}
```

The merged loop is equivalent to the one above, and has no loop-carried dependencies either. So it can be parallelized.

Difficult dependencies


```
for (i=1; i<n; i++) {  
  for (j=1; j<m; j++) {  
    a[i][j] = a[i-1][j-1] +  
              a[i-1][j] +  
              a[i][j-1];  
  }  
}
```



Difficult dependences

- It is not possible to parallelize the loop iterations no matter what loop we consider


```
for (i=1; i<n; i++) {  
  for (j=1; j<m; j++) {  
    a[i][j] = a[i-1][j-1] +  
              a[i-1][j] +  
              a[i][j-1];  
  }  
}
```



Difficult dependences


- It is not possible to parallelize the loop iterations no matter what loop we consider

```
for (i=1; i<n; i++) {  
  for (j=1; j<m; j++) {  
    a[i][j] = a[i-1][j-1] +  
              a[i-1][j] +  
              a[i][j-1];  
  }  
}
```



Solution

- It is possible to parallelize the inner loop by sweeping the matrix diagonally
 - Wavefront sweep*



```
for (slice=0; slice < n + m - 1; slice++) {  
    z1 = slice < m ? 0 : slice - m + 1;  
    z2 = slice < n ? 0 : slice - n + 1;  
    /* The following loop can be parallelized */  
    for (i = slice - z2; i >= z1; i--) {  
        j = slice - i;  
        /* process a[i][j] ... */  
    }  
}
```

Exercises

Exercise

- Can this loop be parallelized? How?

```
#define N some_big_number
double a[N];
const double f = ...;
double val = 1.0;
int i;

for (i=0; i<N; i++) {
    a[i] = val;
    val = val * f
}
```

Exercise

- Which loop(s) (if any) can be parallelized? Why?

```
#define N some_big_number
int s[N] = {1, 1, ... 1};
double p[N] = { ... };

/* Assume that function f() has no side effects */

for (int i=0; i<N; i++) {
    if (s[i]) {
        for (int j=0; j<N; j++) {
            if (s[j] && f(p[i], p[j])) {
                s[j] = 0;
            }
        }
    }
}
```

Exercise

- Which loop(s) can be parallelized? Why?

```
#define N 10000
double phi[2][N][N], maxdelta;
const double EPS = 1.0e-6;
int cur = 0, next = 1;

/* ...Initializations not shown... */

do {
    maxdelta = 0.0;
    for (int i=1; i<N-1; i++) {
        for (int j=1; j<N-1; j++) {
            phi[next][i][j] = (phi[cur][i+1][j] + phi[cur][i-1][j] +
                               phi[cur][i][j+1] + phi[cur][i][j-1]) / 4;
            const double delta = fabs(phi[next][i][j] - phi[cur][i][j]);
            if (delta > maxdelta) {
                maxdelta = delta;
            }
        }
    }
    /* exchange "cur" and "next" */
    const int tmp = cur; cur = next; next = tmp;
} while (maxdelta > EPS);
```

Conclusions

- A loop can be parallelized if there are no dependences crossing loop iterations
- Some kinds of dependences can be eliminated by
 - applying code transformations (*reordering*, *aligning*)
 - using patterns (e.g., *scan*, *reduction*)
 - sweeping across iterations "diagonally"
- There are cases where dependences can simply not be removed